

SNIFF

The AI mystery shopper for your website

An autonomous browser agent that tests signup flows and audits landing pages the way a real, easily-confused human would — then tells you exactly why they left.

sniff-web.vercel.app

Hackathon submission · 2026

THE PROBLEM

“Why aren’t people signing up?” has no honest, evidence-backed answer today.

6.6%

median landing-page
conversion rate across all
industries

*Unbounce, 2025 Conversion Benchmark
Report*

70.2%

average documented
checkout/signup
abandonment rate

*Baymard Institute, 2025 (50-study
meta-analysis)*

\$260B

in recoverable lost orders
across US+EU e-
commerce from fixable
checkout/UX friction

*Baymard Institute, checkout usability
research*

THE SOLUTION

An agent that behaves like a real visitor, not a script.

TRADITIONAL QA

- Hardcoded CSS selectors
- Breaks the moment the UI changes
- Tells you pass/fail, never why
- Can't judge confusing copy or clutter

SNIFF

- Looks at the real screenshot, like a person
- Decides the next click from what it sees
- Diagnoses root cause, severity, owner
- Scores clarity, trust, and friction

directly

THE PRODUCT

Two surfaces, one agent.

01

Signup / Onboarding Testing

- Persona-driven runs (confused first-timer, impatient user, careful reader)
- Root-cause diagnosis: Backend / UX / Performance / Integration
- Severity classification P0–P3, escalated via Slack

02

Landing-Page Conversion Audit

- Overall score + 5-dimension breakdown (clarity, fit, action path, trust, depth)
- Annotated screenshot of exactly what it clicked
- Step-by-step narrative, copy rewrites, prioritized fixes — nothing paywalled

HOW IT WORKS

Three-tier reasoning stack — smart where it matters, cheap everywhere else.

TIER 1 · DETERMINISTIC

Free DOM/CSS checks

Colors, fonts, link health, Core Web Vitals, SEO/meta — zero model calls

TIER 2 · JEV (ULTRA-FAST)

Sanity gate & routing

Goal-reached checks, context enrichment, decision critique

TIER 3 · GEMINI + K2-HORIZON

Vision-grounded judgment


Navigation decisions, scoring, narrative, copy rewrites

MARKET

The market this sits in is compounding fast.

AI Test Automation Market
(\$B, 22.3% CAGR)

AI Agents Market
(\$B, 49.6% CAGR)




Sources: MarketsandMarkets, "AI Test Automation Market" (2026) — [marketsandmarkets.com](https://www.marketsandmarkets.com)

Grand View Research, "AI Agents Market Size To Reach \$182.97 Billion By 2033" — [grandviewresearch.com](https://www.grandviewresearch.com)

WHY NOW

Enterprises are already moving — most just haven't scaled yet.



Gartner forecasts 70% of enterprises will have AI-augmented software testing tools in their engineering toolchain by 2028, up from ~20% in early 2025.

Sources: McKinsey & Company, "The State of AI in 2025" (2025) — [mckinsey.com](https://www.mckinsey.com)

Capgemini, "World Quality Report 2025–26" (17th ed.) — [capgemini.com](https://www.capgemini.com) · Gartner, AI-Augmented Software Testing Tools Magic Quadrant (Oct 2025)

TRACTION

Built, tested, and live — not a mockup.

153 / 153

automated tests passing

3

AI providers wired live
— Gemini, k2-horizon,
Jev

2

products in one agent
— signup testing +
conversion audit

100%

public access — no
signup required for
judges

Deployed end-to-end: FastAPI backend on Cloud Run, Next.js frontend on Vercel. Verified with real runs and real audits against the live stack.

Try it live.

Paste any URL and watch an AI agent tell you the truth about your website.

sniff-web.vercel.app